

KMT-2026-BLG-0521 / OGLE-2026-BLG-0058

Astrometric analysis — summary

Prepared 2026-09-06. Solution: `~/KMTdata/Results/v3_Final/IFfinal_260058_CTIO_.mat` (variable IFsys). Code at commit `47c52cf` on branch `dev1`.

1. Data

quantity	BLG41	BLG01
exposures	19 133	19 981
after cuts	17 354	17 687
sources fitted	594	621
nights	1 985	2 009

KMT CTIO, 2016–2026, I band. Two **overlapping cut-outs** of the same star, 300×300 pix at $0.4''/\text{pix}$ ($2' \times 2'$). They share **no exposures**: the two fields are observed alternately, a median of 12.2 min apart, on 1 947 common nights, so they are independent measurements.

V-band frames exist for 2016–2022 ($\sim 1\,340$ per field) and are used only to check the colour; they are not used in the astrometry.

2. What is in the fit

Per epoch

- full affine transformation, 6 parameters: translation, rotation, scale, shear

Per source

- position (x_0, y_0) and proper motion (μ_x, μ_y) — 4 parameters

Detrending terms

- **differential chromatic refraction (DCR)** — `[1, sin(pa)·secz, cos(pa)·secz]` plus higher orders in `secz·sin(pa)` and `secz·cos(pa)`, fitted in **6 colour bins** (equal population). *This is included.*
- **annual term**
- **pixel-phase correction**, from the per-epoch registration shifts
- **SysRem**, 2 components on the astrometric residuals, computed **once over the whole decade** and reused by every sub-fit
- iterative reweighting from the residuals, with moving-median outlier rejection

Convergence. 15 weighted iterations before SysRem, 6 after. The last three iterations move the bright-star median by 0.003 mas (BLG41) and 0.001 (BLG01).

3. What is not in the fit

- **parallax** — bulge sources give ≈ 0.12 mas, far below the noise floor
 - **any microlensing or astrometric-deviation model** — the solution is model-free; no theoretical curve of any kind has been fitted or subtracted
 - **absolute astrometry** — there is no Gaia tie. Proper motions are *relative*, in the registered pixel frame; 400 mas/pix is assumed, not measured, and the orientation on the sky is unknown
 - acceleration or binary motion
 - any correction for the faint-end photometric bias (§6)
-

4. Steps

1. **KMT_pipelineI** — registration, PSF photometry, source matching → MatchedSources; cross-match to the OGLE catalogue for I and V, using a **per-field** pixel offset (BLG41 229.202/229.283, BLG01 227.542/229.689). 2. **ml.util.mmsFromMatchedSources** — magnitudes rebuilt from FLUX_PSF (the pipeline leaves every MAG_ column empty), per-epoch zero point anchored on OGLE I, then SysRem photometry; airmass and parallactic angle computed from JD and the field centre; colour from OGLE V–I; quality cuts. 3. **Joint fit 2016–2025** → proper motions. The event season is excluded so that it cannot help define the motion it is later measured against. 4. **Full-decade fit** with those motions held → one SysRem correction for the run. 5. **Ten single-season fits** reusing that correction. 6. **Final global fit, 2016–2026, motions solved**. This is the solution behind every plot in this report.*

5. Calibrating stars — selection, and their colour in the plots

Selection (all conditions required):

- $14 < I_{\text{OGLE}} < 19$
- **no companion brighter than $I = 18$ within 2.0 arcsec**, i.e. **5 KMT pixels** at 0.4"/pix. The companion list is the OGLE catalogue, which is at 0.26"/pix and resolves about six times more sources than KMT detects, mapped into KMT pixels with the per-field offset above.
- a star's own OGLE counterpart is excluded from its companion test.

Surviving: **287 stars in BLG41, 253 in BLG01**.

Colour key in the RMS plots

colour	meaning
grey dots	all field stars
blue circles	calibrating stars , as selected above
black line	running median of <i>all</i> stars, in 0.5 mag bins
red star	the target

One point to be clear about: in this solution the per-epoch transformation is fitted from **all** sources, each weighted by its own residual scatter. The calibrating stars are a clean, well-measured subset shown for reference; they do not exclusively define the frame. (UseRefSources, which would restrict the frame fit to them, is available but is **off** here.)

6. Caveats that matter for reading the plots

The measured magnitudes run faint at the faint end. KMT - I_OGLE is flat to ± 0.04 down to $I = 17$, then rises to **+0.50 by $I = 18.75$** . The target's own measured magnitude is 18.70 against its catalogue value of **18.13**. All plots here therefore use the **OGLE catalogue I** on the magnitude axis. PSF and aperture photometry bracket OGLE and disagree by ~ 0.96 mag at $I = 18.5$, which points to crowding: OGLE lists 3 783 sources in this cut-out where KMT detects 594. **The astrometry is not affected** — positions come from the PSF centroid, not the flux, and removing the photometric anchor entirely was measured to cost 0.012 mas.

The CSV carries two error options; choose deliberately.

column	what it is
errX_decade, errY_decade	the black running-median curve of the RMS plot read at the target's catalogue magnitude $I = 18.13$. Constant.
errX_season, errY_season	the target's own residual RMS within its observing season. Follows its brightness.

Neither is a per-epoch measurement uncertainty. The decade value is a population estimate at a fixed magnitude; the season value is the target's own scatter.

They differ by a factor of two where it matters most. The target is at $I \approx 18.1$ for nine seasons and is magnified by about 1.1 mag in 2026, so:

season	target mag	errX / errY, BLG41	BLG01
1 (2016)	18.73	19.28 / 18.22	19.32 / 19.51
4 (2019)	18.71	31.92 / 26.31	29.21 / 26.56
6 (2022)	18.74	33.81 / 27.17	38.55 / 31.45
9 (2025)	18.61	16.85 / 15.13	16.76 / 14.75
10 (2026)	17.65	12.46 / 11.38	11.82 / 11.54
decade constant	—	21.80 / 23.48	22.68 / 22.63

Using the constant would overstate the uncertainty during the event by about a factor two, and understate it in 2022 by a third.

Plotted positions are relative to the target's own mean position, not to the field centre and not to any absolute reference. Each motion curve has the target's decade-average position subtracted, so the zero line is that average. The frame itself is set by the ensemble of 594 (BLG41) and 621 (BLG01) field stars through the per-epoch transformation, not by the geometric centre of the cut-out: differences along a curve are meaningful, the absolute level is not, and the level cannot be compared between the two fields, whose cut-outs have different origins and independently fitted frames. For scale, the target lies +0.42, +0.39 pix (+168, +155 mas) from the centre of the BLG41 cut-out and +1.55, +0.15 pix (+620, +60 mas) from the centre of BLG01; the centring absorbs both.

Only sources with an OGLE counterpart appear in the RMS plots (540 of 594 in BLG41, 478 of 621 in BLG01), which biases the plotted sample slightly bright.

7. Where the numbers stand

quantity	BLG41	BLG01
bright-star residual ($I < 17$), $\Delta X/\Delta Y$	6.414 / 6.879 mas	6.475 / 7.007 mas
calibrating stars, $\Delta X/\Delta Y$	8.29 / 8.30	7.47 / 8.14
target, whole decade	24.62 / 21.00	25.48 / 21.78
target, 2026 season only	$\approx 11.7 / 11.1$	$\approx 12.2 / 11.7$

The target's decade figure is dominated by the nine pre-event seasons where it sits at $I \approx 18.1$; in 2026 the event brightens it by about 1.1 mag and its residual halves. **12 mas is the target's 2026 precision, 24 mas its decade-average precision** — they are different quantities.

8. Files in this report

file	contents
report_RMS_afterPM.png	residual RMS against OGLE I, per axis, per field — after position and proper motion removed
report_motion___nobin.png	source motion, unbinned; top = position with PM retained , bottom = residual
report_motion___binsid.png	the same, binned in one sidereal month (27.321661 d); error bars are the RMS within the bin
source_motion_.csv	JD, X_mas, Y_mas, errX_decade, errY_decade, errX_season, errY_season, season

Eight motion figures in total: 2 fields $\times$ 2 axes $\times$ 2 binnings.

Figures




Figure 1. Residual RMS against the OGLE catalogue magnitude, per axis and per field, **after position and proper motion have been removed**. Grey: all field stars. Blue: calibrating stars ($14 < I < 19$, no $I < 18$ companion within 2 arcsec). Black: running median of all stars, the curve the CSV decade errors are read from. Red: the target.

report_RMS_afterPM.png




Figure 2. **BLG41, X** — unbinned, one point per exposure, no error bars. Top: the measured X position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.
report_motion_BLG41_X_nobin.png




Figure 3. **BLG41, X** — binned in one sidereal month (27.321661 d); error bars are the RMS within the bin. Top: the measured X position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG41_X_binsid.png




Figure 4. **BLG41, Y** — unbinned, one point per exposure, no error bars. Top: the measured Y position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG41_Y_nobin.png




Figure 5. **BLG41, Y** — binned in one sidereal month (27.321661 d); error bars are the RMS within the bin. Top: the measured Y position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG41_Y_binsid.png




Figure 6. **BLG01, X** — unbinned, one point per exposure, no error bars. Top: the measured X position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG01_X_nobin.png




Figure 7. **BLG01, X** — binned in one sidereal month (27.321661 d); error bars are the RMS within the bin. Top: the measured X position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG01_X_binsid.png




Figure 8. **BLG01, Y** — unbinned, one point per exposure, no error bars. Top: the measured Y position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG01_Y_nobin.png




Figure 9. **BLG01, Y** — binned in one sidereal month (27.321661 d); error bars are the RMS within the bin. Top: the measured Y position with the proper motion **retained**; the red line is the fitted proper motion. Bottom: the same with position and proper motion removed. Positions are relative to the target's own mean. The vertical range is set from the robust scatter, so a handful of extreme points fall outside it.

report_motion_BLG01_Y_binsid.png