

KMT-2026-BLG-0521 / OGLE-2026-BLG-0058

Astrometric analysis — summary

Prepared 2026-09-06. Solution: ~/KMTdata/Results/v3_Final/IFfinal_260058_CTIO_<field>.mat

1. Data

quantity	BLG41	BLG01
exposures	19 133	19 981
after cuts	17 354	17 687
sources fitted	594	621
nights	1 985	2 009

KMT CTIO, 2016–2026, I band. Two **overlapping cut-outs** of the same star, 300 × 300 pix at 0.4"/pix (2' × 2'). They share **no exposures**: the two fields are observed alternately, a median of 12.2 min apart, on 1 947 common nights, so they are independent measurements.

V-band frames exist for 2016–2022 (~1 340 per field) and are used only to check the colour; they are not used in the astrometry.

2. What is in the fit

Per epoch

- full affine transformation, 6 parameters: translation, rotation, scale, shear

Per source

- position (x_0 , y_0) and proper motion (μ_x , μ_y) — 4 parameters

Detrending terms

- **differential chromatic refraction (DCR)** — $[1, \sin(pa) \cdot \text{secz}, \cos(pa) \cdot \text{secz}]$ plus higher orders in $\text{secz} \cdot \sin(pa)$ and $\text{secz} \cdot \cos(pa)$, fitted in **6 colour bins** (equal population).
- **annual term**
- **pixel-phase correction**, from the per-epoch registration shifts
- **SysRem**, 2 components on the astrometric residuals, computed **once over the whole decade** and reused by every sub-fit
- iterative reweighting from the residuals, with moving-median outlier rejection

Convergence. 15 weighted iterations before SysRem, 6 after. The last three iterations move the bright-star median by 0.003 mas (BLG41) and 0.001 (BLG01).

3. What is not in the fit

- **parallax** — bulge sources give ≈ 0.12 mas, far below the noise floor
- **any microlensing or astrometric-deviation model** — the solution is model-free; no theoretical curve of any kind has been fitted or subtracted
- **absolute astrometry inside the fit** — the solution itself is *relative*: proper motions are in the registered pixel frame, 400 mas/pix is assumed rather than measured, and the orientation is not used. A Gaia tie is available afterwards (`ml.util.gaiaTie`, see section 7) and gives the orientation, the scale and absolute proper motions, but it is applied to the finished solution and does not constrain the fit
- acceleration or binary motion
- any correction for the faint-end photometric bias (§6)

4. Steps

1. **KMT_pipelineI** — registration, PSF photometry, source matching → `MatchedSources`; cross-match to the OGLE catalogue for I and V, using a **per-field** pixel offset (BLG41 229.202/229.283, BLG01 227.542/229.689). 2. **ml.util.mmsFromMatchedSources** — magnitudes rebuilt from `FLUX_PSF` (the pipeline leaves every `MAG_*` column empty), per-epoch zero point anchored on OGLE I, then `SysRem` photometry; airmass and parallactic angle computed from JD and the field centre; colour from OGLE V–I; quality cuts. 3. **Joint fit 2016–2025** → proper motions 4. **Full-decade fit** with those motions held → one `SysRem` correction for the run. 5. **Ten single-season fits** reusing that correction. 6. **Final global fit, 2016–2026, motions solved.**

5. Calibrating stars — selection, and their colour in the plots

Selection (all conditions required):

- $14 < I_{\text{OGLE}} < 19$
- **no companion brighter than $I = 18$ within 2.0 arcsec**, i.e. **5 KMT pixels** at $0.4''/\text{pix}$. The companion list is the OGLE catalogue, which is at $0.26''/\text{pix}$ and resolves about six times more sources than KMT detects, mapped into KMT pixels with the per-field offset above.
- a star’s own OGLE counterpart is excluded from its companion test.

Surviving: **287 stars in BLG41, 253 in BLG01** — these are the blue points in Figure 1, and are *not* the smaller set the reference-frame test uses. See the table below.

Colour key in the RMS plots

colour	meaning
grey dots	all field stars
blue circles	calibrating stars , as selected above
black line	running median of <i>all</i> stars, in 0.5 mag bins
red star	the target

Two different selections appear in this report, and they must not be confused. Both use the isolation rule above; they differ only in the magnitude window:

	window	in window	isolated	share of all sources
calibrating stars — the blue points in Figure 1	14 < I < 19	529 / 470	287 / 253	48% / 41%
frame stars — what UseRefSources would fit the frame from	14 < I < 16	87 / 62	36 / 35	6% / 6%

BLG41 first, BLG01 second. The blue points are "clean" only in the sense of being un-blended: the 251 extra stars on BLG41 are the faint half, $16 < I < 19$, and they measure 9.06 / 9.51 mas, essentially the field average of 9.90 / 10.23. The 36 bright ones measure 3.89 / 3.79. Requiring brightness as well as isolation is what cuts 287 down to 36.

One point to be clear about: in this solution the per-epoch transformation is fitted from **all** sources, each weighted by its own residual scatter. Neither of the two sets above defines the frame. UseRefSources, which would restrict the frame fit to the 36 / 35 frame stars — not to the 287 / 253 blue points — is available but is **off** here.

Why the frame is fitted from all stars. This was measured rather than assumed. UseRefSources was run at full scale on both fields, selecting $14 < I < 16$ and no companion brighter than $I = 18$ within a given radius — the frame-star row of the table above, not the blue points — everything else at the defaults:

run	ref stars	rstd bright X/Y	seasonal wander
BLG41, frame from all sources	594	8.85 / 9.23	2.21 / 2.57
BLG41, companion radius 5 pix	36	9.06 / 9.73	1.97 / 2.80
BLG41, companion radius 3 pix	48	8.67 / 9.91	1.96 / 2.82
BLG01, frame from all sources	621	6.85 / 7.40	~1.5 / 1.3
BLG01, companion radius 5 pix	10	17.27 / 14.55	5.74 / 3.55

mas. **BLG01 fails outright** — two and a half times worse, with the field's seasonal wander nearly quadrupled. **BLG41 is roughly neutral**, not clearly worse: the bright rstd degrades slightly while the event source's seasonal wander improves by 11% in X.

That BLG01 run was compromised, and has now been redone. The test ran on 2026-08-26; the per-field OGLE registration fix (e89b484) landed on 2026-09-03. Until then BLG41's offset was applied to BLG01, holding its OGLE coverage at 23.5% instead of 60.9% — and the reference selection needs an OGLE magnitude, so it drew from a starved pool: **10 isolated stars, against 35 with the offsets right**. BLG41, which the bug never touched, had 36 both times.

The whole four-step chain was therefore rerun on both fields with UseRefSources and the correctly registered catalogue, against the same chain fitting the frame from all sources. Identical inputs, identical convergence settings:

field	frame	bright rstd X/Y	target rstd X/Y	bright wander	target wander
BLG41	all 594	6.414 / 6.879	24.62 / 21.00	2.62 / 2.77	10.75 / 7.00
BLG41	36 refs	6.569 / 6.905	24.54 / 20.83	2.69 / 2.97	11.49 / 7.03
BLG01	all 621	6.475 / 7.007	25.48 / 21.78	2.83 / 3.06	13.63 / 9.13
BLG01	35 refs	6.445 / 7.225	25.21 / 21.84	2.84 / 3.01	13.66 / 9.00

mas. **BLG01's failure is gone.** With 35 well-distributed reference stars in place of 10 it now behaves exactly as BLG41

always did: every difference is within 3%, and it falls both ways. The 2.5x collapse was an artefact of the registration bug, not a property of the field or of the method.

Neither field gains anything. Across both fields the reference frame moves the bright residual by at most 3%, improves the target marginally on BLG41 and BLG01 in X, and degrades it slightly in Y. Nothing here approaches the factor of three to five the target would need.

The tie to Gaia says the same. Re-tying the reference-frame solutions, the target's **absolute** proper motion is stable against the frame choice even though the relative one is not:

field	relative, all -> refs	absolute, all -> refs
BLG41	+2.20 / -2.31 -> +2.50 / -1.82	-2.92 / -7.25 -> -3.16 / -7.18
BLG01	-0.66 / -2.90 -> -0.45 / -2.94	-2.22 / -7.24 -> -2.25 / -7.29

mas/yr. The relative motion shifts by up to 0.49 mas/yr, the absolute by at most 0.24 — confirming that the shift is gauge, as restricting the frame to a different set of stars must produce. But the all-source solutions tie *better*: their two fields' rotations agree to **0.22 deg** (131.01 and 130.79) where the reference-frame ones differ by **3.8 deg** (133.24 and 129.42), and their per-source scatter about the Gaia relation is lower (3.75/3.83 against 4.54/4.29 on BLG41). That is a further, independent argument for the full field.

Why the subset cannot win here. The reference stars are individually excellent — 36 survivors on BLG41 measure 3.89 / 3.79 mas and are detected in **100%** of epochs, against 9.90 / 10.23 mas and 81% for the field — but they are few, and two numbers show the cost:

	ideal frame noise	share of field inverse variance	leverage at target	max leverage
BLG41, all 594 sources	0.295 / 0.310 mas	100%	0.0017	0.011
BLG41, 36 references	0.653 / 0.673 mas	20%	0.0290	0.248
BLG01, all 621 sources	0.297 / 0.312 mas	100%	0.0016	0.010
BLG01, 35 references	0.662 / 0.700 mas	20%	0.0337	0.416
BLG01, 10 references, the bug	—	—	0.0761	0.939

Leverage is $h(r) = a' (\sum a a')^{-1} a$ with $a = [x, y, 1]$: the variance of the frame shift predicted at a field position, in units of one reference star's variance. The bugged selection reaches a max leverage near 1 — the frame at the field corner no better determined than a single star — with its stars spanning only 55 x 72 pix against 92 x 88 for the full field. That is the mechanism by which ten clustered stars wrecked the solution, though the precise amplification to a factor 2.5 was never quantified and, the run having been redone, no longer needs to be.

Two further reasons, independent of counting:

- The fit already does this, more gently.** Each source is weighted by its own residual scatter — `calculateNee` takes `median(W,1,'omitnan')` per source — so a noisy faint star already contributes almost nothing. `RefSrcFlag` replaces that soft weighting with a hard cut. The effective number of sources behind the weighted frame is 223 (BLG41) and 225 (BLG01), not 594 and 621: the fit has already done most of the selecting.

- **The dominant noise is shared.** The two-field comparison put about 74% of the nightly noise in the atmosphere, common to every star. Reference stars sit under the same atmosphere as the target, so a cleaner frame cannot reach it. Only the ~26% reduction noise is attackable, and 26% in variance is 15% in amplitude, which caps the possible gain — about the size of the shifts in the table above.

The approach is sound and remains implemented. UseRefSources stays **off** by default, now on a clean measurement of both fields rather than one clean field and one corrupted run.

6. The event is not confined to one season

With the OGLE parameters $t_0 = \text{JD } 2461214.2$ (2026-06-22), $t_E = 231.7$ d and $u_0 = 0.499$, the magnification at each season's mean epoch is as below.

- **u** is the angular separation between lens and source at that epoch, in units of the Einstein radius: $u = \sqrt{u_0^2 + ((t - t_0)/t_E)^2}$. It is dimensionless, falls to $u_0 = 0.499$ at closest approach, and is large long before and after the event.
- **A** is the **flux** magnification, $A = (u^2 + 2) / (u \cdot \sqrt{u^2 + 4})$. $A = 1$ means unmagnified; $A = 2$ means twice as much light.
- **Δmag** is the same quantity in magnitudes, $\Delta\text{mag} = -2.5 \cdot \log_{10}(A)$. The two columns are not independent: $A = 2.067$ is $\Delta\text{mag} = -0.788$, and -0.788 mag inverts to $A = 2.066$. Where the text below quotes "a magnification of about 2.1" and "0.8 mag" it is quoting the same number twice.

season	$(t - t_0)$ [d]	u	A	Δmag
2016	−3644	15.74	1.000	0.000
2019	−2564	11.08	1.000	0.000
2022	−1474	6.38	1.001	−0.001
2023	−1114	4.83	1.003	−0.003
2024	−754	3.29	1.012	−0.013
2025	−394	1.77	1.085	−0.089
2026	−44	0.53	2.067	−0.788

The event is above 1% magnification for **4.4 years**, above 5% for 2.6 years and above 10% for 2.0 years.

The model predicts 0.80 mag of brightening at the 2026 season median, whereas the measured magnitudes give 1.085 (18.724 at baseline against 17.639 in 2026). The 0.29 mag difference is the faint-end photometric bias of section 7, which is +0.39 mag at $l \approx 18.1$ and +0.13 at $l \approx 17.3$: a *difference* of two measured magnitudes inflates by the difference of their biases, $0.80 + 0.26 = 1.06$ against 1.085 observed. It is a consistency check on the bias, not a discrepancy with the model.

Calling 2026 "the event season" and 2016–2025 "pre-event" is therefore wrong: **2025 is already magnified by 9%** and 2024 by 1.3%.

This matters for the proper-motion baseline. The astrometric deviation of a microlensing event does not peak with the photometric one: for these parameters it is largest near $u = \sqrt{2}$, which falls in **2025** (2.46 mas), with 2024 at 1.91 mas and 2026 back down to 1.63 mas as the source passes closest approach. So the joint fit that supplies the proper motions, which holds out 2026 alone, still contains the two seasons carrying the **largest** astrometric signal.

Any proper motion quoted here is therefore fitted through the event, and would absorb part of a real astrometric excursion. A baseline free of it would have to stop around 2022, which costs a third of the time span and most of the leverage on the proper motion. Nothing in this report depends on that choice — no astrometric model is fitted — but a future attempt to measure the deviation must not treat 2016–2025 as an event-free baseline.

7. Caveats that matter for reading the plots

The measured magnitudes run faint at the faint end. `KMT - I_OGLE` is flat to ± 0.04 down to $I = 17$, then rises to **+0.50 by $I = 18.75$** . The target's own measured magnitude is 18.70 against its catalogue value of **18.13**. All plots here therefore use the **OGLE catalogue I** on the magnitude axis. PSF and aperture photometry bracket OGLE and disagree by ~ 0.96 mag at $I = 18.5$, which points to crowding: OGLE lists 3 783 sources in this cut-out where KMT detects 594. **The astrometry is not affected** — positions come from the PSF centroid, not the flux, and removing the photometric anchor entirely was measured to cost 0.012 mas.

The CSV carries two error options; choose deliberately.

column	what it is
<code>errX_decade</code> , <code>errY_decade</code>	the black running-median curve of the RMS plot read at the target's catalogue magnitude $I = 18.13$. Constant.
<code>errX_season</code> , <code>errY_season</code>	the target's own residual RMS within its observing season. Follows its brightness.

Neither is a per-epoch measurement uncertainty. The decade value is a population estimate at a fixed magnitude; the season value is the target's own scatter.

They differ by a factor of two where it matters most. The target is at $I \approx 18.1$ for nine seasons and is magnified by 0.80 mag in 2026, so:

season	target mag	<code>errX / errY</code> , BLG41	BLG01
1 (2016)	18.73	19.28 / 18.22	19.32 / 19.51
4 (2019)	18.71	31.92 / 26.31	29.21 / 26.56
6 (2022)	18.74	33.81 / 27.17	38.55 / 31.45
9 (2025)	18.61	16.85 / 15.13	16.76 / 14.75
10 (2026)	17.65	12.46 / 11.38	11.82 / 11.54
decade constant	—	21.80 / 23.48	22.68 / 22.63

Using the constant would overstate the uncertainty at the peak of the event by about a factor two, and understate it in 2022 by a third.

The sign of a proper-motion component carries no physical meaning. The fitted motion of the target is +2.20 / -2.31 mas/yr in BLG41 and -0.66 / -2.90 in BLG01: the same star, the same nights, opposite signs in X. This is a gauge difference between the two frames, not anything about the source. The per-epoch transformation is free at every epoch, so it absorbs any motion common to the field, and what survives is defined only relative to the mean motion of each cut-out's own star population. Those populations differ:

quantity	BLG41	BLG01
median PM of the field's own sources	+2.128 / +0.183	-0.261 / -0.617

Measured directly on the 383 well-measured stars common to both fields, the disagreement is the same for every star:

quantity	X	Y
PM difference, BLG01 – BLG41, all common stars	-2.41 ± 0.43	-0.60 ± 0.57
the target	-2.86	-0.59

The target's offset is the population's offset. Fitting the difference as a constant plus a linear term in position gives, in addition, a relative rotation of 0.28 deg/yr and a scale drift of $4.7\text{e-}3$ /yr between the two frames; removing both leaves 1.14 / 0.35 mas/yr, consistent with per-star measurement error.

So the Y values (–2.31 and –2.90) agree only because the Y gauge offset is small, and the X values disagree only because the X offset is large. **What is physical is the differential motion between stars within one field**, unless the frame is tied to an external one — which it now can be.

Tied to Gaia, the two fields agree. `ml.util.gaiaTie` carries a solution onto the Gaia frame through OGLE, giving the orientation, the plate scale and an absolute proper motion for every matched source:

quantity	BLG41	BLG01
KMT sources matched to Gaia	532 of 594 (90%)	447 of 621 (72%)
median separation	0.259 arcsec	0.358 arcsec
calibrators (Gaia PM error < 0.4 mas/yr, $l < 18$)	426	327
fitted rotation of the pixel axes on the sky	+131.01 deg	+130.79 deg
gauge offset removed	$-0.36 / +4.41$ mas/yr	$-2.84 / +3.88$
scatter about the Gaia relation	$3.75 / 3.83$ mas/yr	$3.38 / 3.51$

The target's proper motion:

frame	BLG41	BLG01	difference
relative (detector x, y), as fitted	$+2.196 / -2.310$	$-0.660 / -2.900$	$2.856 / 0.590$
absolute ($\mu_\alpha \cos \delta$, μ_δ), Gaia frame	$-2.917 / -7.248$	$-2.215 / -7.244$	$0.702 / 0.004$

mas/yr, taken from `IF.ParS(3:4,Ie)` at 400 mas/pix, target index 383 (BLG41) and 450 (BLG01). **The two rows are not in the same axes:** the relative values are in the cut-out's detector x/y, the absolute ones are ($\mu_\alpha \cos \delta$, μ_δ), since `gaiaTie` returns `Tie.A \ (PM - Offset)` and so carries the pixel frame into the Gaia frame through a 131 deg rotation with parity and a 0.92 scale.

The sign disagreement disappears once both are on the same physical frame, and the two rotations, derived independently from separately reduced data, agree to 0.22 deg — which is what makes the tie credible rather than a fit to noise. The relative disagreement itself is gauge, not astrophysics: both cut-outs have nearly the same orientation (131.01 and 130.79 deg), so the +2.20 against –0.66 in x is the unconstrained constant-plus-gradient freedom of each field's own frame, which is exactly what the tie removes.

Against Gaia's own value for this star. The target has a Gaia counterpart, so the absolute motion can be checked directly rather than only field against field:

	$\mu_\alpha \cos \delta$	μ_δ
Gaia DR3, the target	-2.479	-9.129

ours, BLG41	−2.917 (−0.44)	−7.248 (+1.88)
ours, BLG01	−2.215 (+0.26)	−7.244 (+1.89)

mas/yr, with the difference from Gaia in brackets. $\mu_\alpha \cos \delta$ agrees well. μ_δ sits 1.88 mas/yr away, and by nearly the same amount in both fields — which should **not** be read as two independent confirmations, because BLG01 reused BLG41's OGLE-to-Gaia solution and a common systematic in that bridge would reproduce this pattern exactly. Against the per-source scatter it is about 0.5 sigma and unremarkable. It is not the event: 1.88 mas/yr over the decade baseline is some 19 mas of accumulated displacement, an order of magnitude beyond the ~2.5 mas peak deviation the model allows.

Which error applies. Three different numbers circulate and they measure different things:

- **3.4 to 3.8 mas/yr** — scatter about the Gaia relation, *per source*. This is the one that applies to a single star's absolute proper motion, the target included.
- **0.16 / 0.17 mas/yr** — the error on the tie *transformation*, from 532 and 447 matched sources. It bounds the frame, not any individual star.
- **~1.0 / 0.7 mas/yr** — the target's own relative proper motion uncertainty, from the scatter of its season means about its fitted line.

Reduction choices contribute far less: the pixel-phase on/off variants give +2.21 / −2.31 and +2.27 / −2.27 (BLG41), −0.65 / −2.89 and −0.71 / −2.94 (BLG01), a spread of about 0.06 mas/yr.

Two cautions. The μ_δ agreement of 0.004 mas/yr between the fields is far better than either field's own precision and is therefore luck; the honest statement is that the two agree to within errors of order 1 mas/yr on the mean. And the tie fixes the **frame, not the precision**: our per-source scatter of 3.4 to 3.8 mas/yr exceeds Gaia's own spread of 2.9 to 3.3, so most of the per-star difference is ours and nothing in the residual analysis improves.

Plotted positions are relative to the target's own mean position, not to the field centre and not to any absolute reference. Each motion curve has the target's decade-average position subtracted, so the zero line is that average. The frame itself is set by the ensemble of 594 (BLG41) and 621 (BLG01) field stars through the per-epoch transformation, not by the geometric centre of the cut-out: differences along a curve are meaningful, the absolute level is not, and the level cannot be compared between the two fields, whose cut-outs have different origins and independently fitted frames. For scale, the target lies +0.42, +0.39 pix (+168, +155 mas) from the centre of the BLG41 cut-out and +1.55, +0.15 pix (+620, +60 mas) from the centre of BLG01; the centring absorbs both.

Only sources with an OGLE counterpart appear in the RMS plots (540 of 594 in BLG41, 478 of 621 in BLG01), which biases the plotted sample slightly bright.

8. Where the numbers stand

quantity	BLG41	BLG01
bright-star residual ($I < 17$), $\Delta X/\Delta Y$	6.414 / 6.879 mas	6.475 / 7.007 mas
calibrating stars, $\Delta X/\Delta Y$	8.29 / 8.30	7.47 / 8.14
target, whole decade	24.62 / 21.00	25.48 / 21.78
target, 2026 season only	≈ 11.7 / 11.1	≈ 12.2 / 11.7

The target's decade figure is dominated by the seasons in which it sits near its baseline $I \approx 18.1$; in 2026 the magnification of about 2.1 brightens it by 0.8 mag and its residual halves. 2025 is brighter than baseline too, by 0.09

mag. **12 mas** is the target's 2026 precision, **24 mas** its decade-average precision — they are different quantities.

9. The Gaia-pinned reduction (v4)

A second, independent reduction was built to a different specification: restrict the analysis to a clean calibration set, and tie the frame to Gaia from the start rather than afterwards. It is reported here in full because it is a fair test of an obvious idea, and because **it did not improve the target**.

What was done.

1. **Calibration set.** $14 < I_{\text{OGLE}} < 17$, no companion brighter than $I = 18$ within 2.0 arcsec, and a Gaia counterpart with **RUWE < 1.4**. The target is added explicitly (at $I = 18.13$ it falls outside the magnitude window). Every other source is dropped at the input, so all measurements and quality cuts act on this set alone: **127 sources on BLG41, 97 on BLG01**. 2. **The frame is pinned to Gaia.** Each calibration star's position *and* proper motion are held at Gaia's values, converted into cut-out pixels, so the per-epoch affine is fitted against Gaia positions propagated to that epoch by Gaia's own motions. 126 and 96 stars are pinned; the target is left free, since pinning it would erase the signal. The six-parameter gauge is gone by construction and no post-hoc tie is needed. 3. **Chain.** With motions supplied by Gaia, the joint proper-motion fit and the per-season fits are unnecessary. Two steps remain: a full-decade SysRem correction, then a final solution reusing it with the pixel-phase term on.

Getting RUWE. It is **not** in the local catsHTM DR3 table, which carries `astrometric_chi2_al` and `astrometric_n_good_obs_al` but neither RUWE nor the `u0(G, BP-RP)` table needed to normalise UWE into it. The true values were queried from the Gaia archive instead (9464 sources in a 3 arcmin cone, cached in `~/KMTdata/GaiaRef/`). This mattered: at $G < 17$, **87%** pass **RUWE < 1.4**, where the unnormalised UWE proxy predicted 74%. Gaia's own proper-motion errors for these stars are **0.070 / 0.044 mas/yr**, some fifty times better than our internal frame.

The Gaia-to-pixel map is fitted directly from the matched pairs, not carried through OGLE: the OGLE bridge is needed only to identify which star is which, and routing coordinates through it inherits OGLE's ~110 mas astrometric error. Direct fitting gives residuals of **23 / 16 mas** (BLG41) and 21 / 20 (BLG01) against 252 / 159 through OGLE, and a scale of 2.519 pix/arcsec against the 2.500 expected at 0.4 arcsec/pix — an independent check that the map is right. The residual is static per star, so it displaces each star's residuals by a constant and does not touch proper motion or any time-dependent signal.

Result: the target is measured worse, not better.

	v3, all sources + post-hoc tie	v4, calibration set + Gaia-pinned
target rstd, BLG41	24.62 / 21.00	29.62 / 22.63
target rstd, BLG01	25.48 / 21.78	28.73 / 23.25
frame-star rstd, BLG41	6.41 / 6.88	6.13 / 5.20
frame-star rstd, BLG01	6.48 / 7.01	5.30 / 5.42

mas. The calibration stars themselves are measured **better** — they are bright and isolated, which is what they were selected for — while the target is **14 to 20% worse**. The explanation is the one already quantified in section 5: a frame built from 126 stars instead of ~600 weighted sources has roughly twice the ideal noise and twenty to fifty times the leverage. Restricting the analysis set is what costs; the Gaia tie is not at fault.

Convergence is not the problem. Over the frame stars the last three iterations move the median by **0.0009 mas (BLG41) and 0.0026 (BLG01)**, comparable to v3.

The two fields do not agree better than before. Converting the target's motion out of pixel axes with each field's own Gaia map:

	BLG41	BLG01	difference
v4, Gaia-pinned	-1.799 / -6.780	-1.158 / -7.147	0.640 / 0.367
v3, post-hoc tie	-2.92 / -7.25	-2.22 / -7.24	0.702 / 0.004

mas/yr. Both agree at about the 1 mas/yr level, which is the target's own proper-motion uncertainty, but pinning the frame did not tighten it. The post-hoc tie was already doing this job.

The comparison stars agree better than the target does. Three stars at the target's own magnitude, outside the calibration set and freely fitted, were carried through both fields as unpinned passengers:

star	BLG41	BLG01	difference
l = 18.14, 4.5 pix from the target	-0.63 / -5.47	-1.08 / -5.47	0.45 / 0.00
l = 17.91, 21 pix	-6.81 / -9.99	-6.94 / -10.06	0.13 / 0.07
l = 18.05, 48 pix	+3.06 / -6.34	+2.80 / -6.48	0.26 / 0.14
the target, l = 18.13	-1.80 / -6.78	-1.16 / -7.15	0.64 / 0.37

mas/yr. Stars of the same brightness, measured the same way in the same frames, reproduce between the two fields two to five times better than the target does. The target's extra disagreement is specific to it, and the obvious candidates are its blending and the fact that its motion is fitted straight through a magnification event lasting several years (section 6).

Sensitivity to the set. Adding just those three passengers moved the fitted target motion by **0.42 / 0.63 mas/yr on BLG41** and 0.09 / 0.41 on BLG01. A result that shifts by half a mas per year when three faint stars join a set of 127 is not a precise result, and this is the most direct measurement in this report of how uncertain the target's motion actually is.

Overlap of the two fields. Matching the two source lists directly, after fitting the offset between the cut-outs (-1.11 / +0.24 pix), gives **546 of 594 BLG41 sources and 575 of 621 BLG01 sources matched within 1 pixel — 92% and 93%**, rising to 97% and 98% at 1.5 pix. These are indeed two observations of nearly the same sky. The unmatched 7% are faint and marginal: median MAG_PSF 17.18 against 16.87 for matched stars, and a detection fraction of 0.54 against 0.84. They are stars one reduction found and the other did not, or blends split differently. Restricted to MAG_PSF < 17 the agreement is 94%.

Matching *through* Gaia and OGLE instead gives a misleading 58%, because their identification coverage over our sources is only 90% / 72% (Gaia) — sources neither catalogue identifies cannot be paired and count as non-overlapping.

Verdict. v4 is a clean, gauge-free reduction whose calibration stars are measured better than anything in v3, and it is the right way to obtain absolute astrometry. It is **not** adopted as the primary solution, because the quantity this project exists to measure — the target's astrometry — comes out worse. The v3 solution with a post-hoc Gaia tie remains the reference.

10. Files in this report

file	contents
------	----------

report_RMS_afterPM.png

residual RMS against OGLE I, per axis, per field — **after position and proper motion**

report_motion_<field>_nobin.png

source motion, unbinned; 2x2 panel, columns X and Y, top row = position with residual

report_motion_<field>_binsid.png

the same, binned in **one sidereal month** (27.321661 d); error bars are the **RMS with**

report_v4_motion_<field>_target.png

v4 motion of the target, sidereal-month bins, 2x2 as above

report_v4_motion_<field>_cal{1,2,3}.png

v4 motion of three calibration stars, proper motion **pinned to Gaia**

report_v4_motion_<field>_out{1,2,3}.png

v4 motion of three I~18 comparison stars outside the set, freely fitted

source_motion_<field>.csv

JD, X_mas, Y_mas, errX_decade, errY_decade, errX_season, errY_season,

Eight motion figures in total: 2 fields × 2 axes × 2 binnings.

Figures




Figure 1. Residual RMS against the OGLE catalogue magnitude, per axis and per field, **after position and proper motion have been removed**. Grey: all field stars. Blue: calibrating stars (14 < I < 19, no I < 18 companion within 2 arcsec). Black: running median of all stars, the curve the CSV decade errors are read from. Red: the target.

report_RMS_afterPM.png




Figure 2. **BLG41** — unbinned, one point per exposure, no error bars. Columns are the two axes, **X left and Y right**. Top row: the measured position with the proper motion **retained**; the red line is the fitted proper motion. Bottom row: the same with position and proper motion removed. All four panels share one vertical scale so the axes can be compared directly; it is set from the robust scatter, so a handful of extreme points fall outside it. Positions are relative to the target's own mean.

report_motion_BLG41_nobin.png




Figure 3. **BLG41** — binned in one sidereal month (27.321661 d); error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**. Top row: the measured position with the proper motion **retained**; the red line is the fitted proper motion. Bottom row: the same with position and proper motion removed. All four panels share one vertical scale so the axes can be compared directly; it is set from the robust scatter, so a handful of extreme points fall outside it. Positions are relative to the target's own mean.

report_motion_BLG41_binsid.png




Figure 4. **BLG01** — unbinned, one point per exposure, no error bars. Columns are the two axes, **X left and Y right**. Top row: the measured position with the proper motion **retained**; the red line is the fitted proper motion. Bottom row: the same with position and proper motion removed. All four panels share one vertical scale so the axes can be compared directly; it is set from the robust scatter, so a handful of extreme points fall outside it. Positions are relative to the target's own mean.

report_motion_BLG01_nobin.png




Figure 5. **BLG01** — binned in one sidereal month (27.321661 d); error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**. Top row: the measured position with the proper motion **retained**; the red line is the fitted proper motion. Bottom row: the same with position and proper motion removed. All four panels share one vertical scale so the axes can be compared directly; it is set from the robust scatter, so a handful of extreme points fall outside it. Positions are relative to the target's own mean.

report_motion_BLG01_binsid.png




Figure 6. **v4, BLG41** — the target, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG41_target.png




Figure 7. **v4, BLG41** — calibration star (pinned to Gaia), faint end of the set. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG41_cal1.png




Figure 8. **v4, BLG41** — calibration star (pinned to Gaia), faint end of the set. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG41_cal2.png




Figure 9. **v4, BLG41** — calibration star (pinned to Gaia), faint end of the set. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG41_cal3.png

BLG41 comparison star outside the set, PM fitted — $I = 18.14$, 5 pix from the target, freely fitted | sky PM ($\mu_\alpha \cos d / \mu_\delta$) $-0.63 / -5.47$ mas/yr




Figure 10. **v4, BLG41** — comparison star outside the set at the target magnitude, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.
report_v4_motion_BLG41_out1.png

BLG41 comparison star outside the set, PM fitted — $I = 17.91$, 21 pix from the target, freely fitted | sky PM ($\mu_\alpha \cos d / \mu_\delta$) $-6.81 / -9.99$ mas/yr




Figure 11. **v4, BLG41** — comparison star outside the set at the target magnitude, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.
report_v4_motion_BLG41_out2.png

BLG41 comparison star outside the set, PM fitted — $I = 18.05$, 48 pix from the target, freely fitted | sky PM ($\mu_\alpha \cos d / \mu_d$) $+3.06 / -6.34$ mas/yr




Figure 12. **v4, BLG41** — comparison star outside the set at the target magnitude, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG41_out3.png




Figure 13. **v4, BLG01** — the target, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG01_target.png




Figure 14. **v4, BLG01** — calibration star (pinned to Gaia), faint end of the set. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG01_cal1.png




Figure 15. **v4, BLG01** — calibration star (pinned to Gaia), faint end of the set. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG01_cal2.png




Figure 16. **v4, BLG01** — calibration star (pinned to Gaia), faint end of the set. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.
 report_v4_motion_BLG01_cal3.png




Figure 17. **v4, BLG01** — comparison star outside the set at the target magnitude, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.
 report_v4_motion_BLG01_out1.png

BLG01 comparison star outside the set, PM fitted — $I = 17.91$, 21 pix from the target, freely fitted | sky PM ($\mu_\alpha \cos d / \mu_\delta$) $-6.94 / -10.06$ mas/yr




Figure 18. **v4, BLG01** — comparison star outside the set at the target magnitude, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG01_out2.png

BLG01 comparison star outside the set, PM fitted — $I = 18.05$, 49 pix from the target, freely fitted | sky PM ($\mu_\alpha \cos d / \mu_\delta$) $+2.80 / -6.48$ mas/yr




Figure 19. **v4, BLG01** — comparison star outside the set at the target magnitude, freely fitted. Binned in one sidereal month; error bars are the RMS within the bin. Columns are the two axes, **X left and Y right**; top row keeps the proper motion, bottom row removes it. The panel heading gives the star's magnitude, its distance from the target and its sky proper motion. **Note that a pinned star's red line is Gaia's proper motion, not a fitted one**, so its bottom panel tests the frame, while the target's tests the fit.

report_v4_motion_BLG01_out3.png