

KMT-2026-BLG-0521 / OGLE-2026-BLG-0058

Astrometry of the microlensing target — summary

Prepared 2026-09-12. Solution: `~/KMTdata/Results/v8/IFfinal_<field>.mat` (variable `IFsys`), tie in `Tie_<field>.mat`.

1. The data and the algorithm

Data. Two KMTNet cut-outs of the same field, **BLG41** and **BLG01**, 300 x 300 pixels at 0.4 arcsec/pix, in I band, spanning 2016 to 2026 in ten observing seasons (2020 is a short stub and is dropped). **19 133** and **19 981** exposures respectively, taken alternately on the same nights, a median of 12.2 minutes apart. The target sits at the field centre at $I_{\text{OGLE}} = 18.13$, RA 17:52:38.09, Dec -31:47:36.1.

Calibration set. Stars are selected on magnitude and isolation: $14 < I < 19$ with **no companion brighter than $I = 18$ within 2 arcsec**, the companion list taken from the OGLE catalogue, which resolves about six times more sources than KMT detects. The set is then **iterated once**: a first solution is computed, and stars whose residual RMS lies more than **2 sigma** above the running median for their magnitude are dropped and the fit repeated. The cut is made in log space, so it is multiplicative and does not simply shave off the faint end. This leaves **260** stars on BLG41 and **235** on BLG01. The target passes the selection on its own merits. Everything else is dropped at the input, so all measurement, detrending and quality cuts act on this set alone.

The fit. Each exposure gets a full 2-D affine transformation, 6 parameters, and each star a position and proper motion, 4 parameters; these are solved together by alternating least squares. Alongside them the fit removes differential chromatic refraction (fitted in 6 equal-population colour bins), an annual term, a pixel-phase term, and two SysRem components computed once over the whole decade. Every star is freely fitted — nothing is held at a catalogue value — and each is weighted by its own residual scatter, with moving-median outlier rejection.

Absolute frame. A free fit determines only *relative* astrometry, so the result is tied to Gaia DR3 afterwards. The tie is fitted on stars with **RUWE < 1.4 and $15 < I < 17$** and removes the full six-parameter gauge — a constant *and* a linear gradient across the field — then is applied to every source.

2. Results

Achieved accuracy, residual RMS over the decade after position and proper motion are removed:

	BLG41 $\Delta X / \Delta Y$	BLG01 $\Delta X / \Delta Y$
calibration stars brighter than $I = 17$	4.80 / 4.83 mas	4.75 / 4.95 mas
the target, $I = 18.13$	23.77 / 19.67 mas	21.51 / 20.45 mas

The frame is therefore good to under **5 mas** on well-measured stars, and the target — three magnitudes fainter and blended — is measured to about **20 mas** per epoch, consistent with the running median of the residual RMS at its magnitude.

Proper motion of the target, absolute, after the Gaia tie:

	$\mu_{\alpha \cos \delta}$	μ_{δ}
BLG41	-1.966 ± 0.347	-7.019 ± 0.358

BLG01

 -1.514 ± 0.227 -7.073 ± 0.376

difference between the fields


 -0.452 ± 0.415 $+0.054 \pm 0.519$

mas/yr. The two independent fields agree within their errors, at 1.09 and 0.10 sigma. The quoted error is internal, from the scatter of the ten season means about the fitted line; the external check against Gaia gives about **1.0 / 0.62 mas/yr** per star, and that is the figure to use in practice. That external scatter is magnitude dependent: 0.55 / 0.35 mas/yr for $I < 17$, 1.8 / 1.0 at 17–18 and 2.5–4.5 / 1.5–2.2 at 18–19, where both KMT blending and Gaia's own errors grow. At the target's magnitude the disagreement with Gaia to expect is therefore 1.5–2 mas/yr in declination, not 0.6.

The blend. The target has a companion of equal brightness (m_{18_d04} , $I = 18.14$, Gaia $G = 19.69$) at 4.5 pixels = 1.8", itself paired with a $G = 20.3$ star 1.0" further out for which Gaia solves no proper motion; under the 3.1" KMT PSF the three are two blobs, and the calibration set's isolation cut, which excludes companions brighter than $I = 18$ within 2", lets the pair through by 0.14 mag. Gaia gives the target -2.48 ± 0.47 / -9.13 ± 0.31 mas/yr, i.e. our declination motion is 2.1 mas/yr less negative than Gaia's in both fields, and the companion $+1.74$ / $+0.27$, 9.4 mas/yr away from the target in declination. Three tests show that this offset is **not** the companion pulling our centroid: the measured separation of the two blobs is 91% of Gaia's, with the shortfall in the direction of the $G = 20.3$ star rather than of the target, so the mutual pull is at most a few per cent and could bias the proper motion by at most ~ 0.5 mas/yr; the position of the target moves with seeing by 5 mas per pixel of FWHM, the same as isolated stars, and the seeing has no trend over the decade; and the relative motion of the two blobs measured directly from the raw positions, without any frame solution, is $+2$ – 3 / $+0.4$ – 0.8 mas/yr against Gaia's $+4.2$ / $+9.4$. The 2.1 mas/yr is inside the 1.5–2 mas/yr KMT-minus-Gaia scatter for stars of this magnitude (26–29% of the stars within 0.5 mag of the target are further from Gaia in declination than it is), and Gaia's own five-parameter solution for the target carries 1.0 mas of astrometric excess noise in a 2" group of three sources. Which of the two declination values is right cannot be decided from these data; what the blend does do for certain is add to the per-epoch scatter of the target, and dilute any astrometric excursion of the source by the flux fraction of the companion inside the fitting aperture.

No astrometric anomaly is detected at the target. Its χ^2 on monthly binned residuals is 1500 and 1313 for 180 degrees of freedom, against field medians of χ^2/DoF of 6.5 and 6.8 — the target sits inside the body of the distribution, not in the tail. Whatever astrometric excursion the microlensing event produces is smaller than the scatter of an ordinary star of this brightness.


Residual RMS after position and proper motion, per axis and field



Residual RMS against the OGLE catalogue magnitude, per axis and per field, after position and proper motion have been removed. Blue: all 260 / 235 calibration stars, every one freely fitted. Green: the stars the Gaia tie is fitted on, RUWE below 1.4 and $15 < I < 17$. Red: the target. Black: running median. The dashed lines mark the 14 and 19 mag selection limits. Compared with a set that has not been iterated, the 2 sigma cut removes the scattered points that previously sat well above the median relation.

report_v8_RMS.png

The target, before and after proper-motion detrending, sidereal-month bins



The target before and after proper-motion detrending, in both fields. Columns are the fields; the top row keeps the proper motion and the bottom row removes it. Blue circles are the offset in right ascension, red squares in declination, binned in one sidereal month with the standard error within the bin. The proper motion is plainly visible as the declination trend in the top row and is absent from the bottom row; the residual RMS quoted in the lower panels is per bin.

report_v8_target_beforeafter.png

More details

3. Object selection, step by step

step	BLG41	BLG01
raw sources in the MSc file	1177	1187
survive the standard quality cuts	594	621
have an OGLE I magnitude	540	478
in $14 < I < 19$	529	470
... and isolated: no $I < 18$ companion within 2.0 arcsec	287	253
... and surviving the 2 sigma RMS cut	260	235

The isolation cut is the dominant one, removing 46% of the magnitude-selected sources on both fields. The outlier iteration then removes a further **27 and 18** stars — 9% and 7% — spread across all magnitudes rather than concentrated at the faint end: 6%, 14%, 10% and 3% by magnitude band on BLG41. Their median residual RMS is 22.4 and 48.5 mas against 10.8 and 10.6 mas for the stars kept.

For the tie only, Gaia is required: 247 of 260 and 207 of 235 sources have a counterpart, of which 198 and 163 pass $\text{RUWE} < 1.4$, and 104 and 85 also fall in $15 < I < 17$.


4. The Gaia tie

The map. Gaia standard coordinates are fitted directly onto cut-out pixels from the tie stars — not carried through OGLE, which would inherit its ~110 mas astrometric error. Median residual **22.1 mas (BLG41)** and **21.8 (BLG01)**, with scales of 2.5195 and 2.5289 pix/arcsec against the 2.500 expected at 0.4"/pix.

The gauge. In the proper motions the gauge freedom is a constant plus a linear gradient across the field — six parameters. Removing only the constant leaves a shear of about -0.036 mas/yr per pixel, reproduced at nearly the same value in both independently reduced fields. With the full gauge removed, the scatter of our proper motions about Gaia is **0.994 / 0.608 mas/yr** (BLG41) and **1.049 / 0.641** (BLG01); removing only the constant would leave 3 to 4 times that. The gauge constant removed is $-0.126 / +5.105$ and $+2.151 / +4.324$ mas/yr.

5. What the fit does, in detail

Each of the two steps runs **two passes** internally. `IFsys.AnnualEffect` and `IFsys.PixPhase` read 0 in the saved solutions, which does *not* mean those effects were ignored: they are fitted and subtracted from the data in pass 1 and switched off in pass 2, which works on data from which they have been removed.



per-source position and motion	fitted	fitted
DCR / chromatic refraction	fitted	fitted again
annual term	fitted and subtracted	forced off
pixel-phase term	fitted and subtracted	forced off
SysRem		applied
iterative reweighting	15 iterations	6 iterations

- **Per-epoch affine** — 6 parameters per exposure in ParE, stored as the deviation from the identity: translation (2), rotation (1), uniform scale (1), shear (2). 520 equations for 6 unknowns each epoch. Translation dominates at about 37 / 28 mas rms; rotation and scale each move a star at the cut-out edge by roughly 10 mas.
- **Per source** — position at JD0 = 2019-06-01 and proper motion, 4 parameters in ParS, in pixels and pixels per year at 400 mas/pix.
- **DCR** — $[1, \sin(\text{pa}) \cdot \text{secz}, \cos(\text{pa}) \cdot \text{secz}]$ plus six higher orders in $(\text{secz} \cdot \sin \text{pa})^n, (\text{secz} \cdot \cos \text{pa})^n$ for $n = 2, 3, 4$: 18 parameters per colour bin, 6 equal-population bins. Only the *differential* part is recoverable — refraction common to the field is exactly a translation and is absorbed by t_x, t_y every epoch.
- **Annual term** — a quartic in the phase of the calendar year, 10 parameters per colour bin, one curve shared by every star of the bin (the same six equal-population colour bins as the DCR term). Not parallax, which is per-source and switched off.
- **Pixel phase** — a quintic in the sub-pixel phase, 10 parameters, global, recoverable only because the pipeline stored the per-epoch registration shifts.
- **SysRem** — 2 components over the whole decade, reused rather than refitted per season, with point-wise rejection of bad cells. Refitting per season was tested and is worse: see section 8.

Timings and convergence:

	BLG41	BLG01
step 1, decade SysRem	16.9 min	15.4 min
step 2, final solution	31.1 min	29.1 min
convergence, last 3 iterations	0.0002 mas	0.0047 mas

No per-season fitting enters the solution; seasons are used only for reporting.

6. Comparison stars

Six, chosen to be the same physical objects in both fields — three near the target's magnitude and three brighter. All six survive the outlier cut and sit at identical positions in the two fields:

tag	I	dist [pix]	decade RMS $\Delta X / \Delta Y$, BLG41	BLG01
m18_d04	18.14	4.5	28.86 / 25.64	25.24 / 27.09
m18_d17	18.09	17.3	44.90 / 23.22	80.36 / 27.53
m18_d48	18.05	48.4	18.06 / 15.57	13.86 / 15.58
m16_d24	16.68	24.4	4.14 / 4.31	4.13 / 4.32

m16_d30	16.63	30.4	6.69 / 5.53	6.55 / 5.62
m16_d31	16.75	30.6	8.42 / 7.82	11.16 / 7.84

mas. The bright trio reproduces between the fields almost exactly — m16_d24 gives $-3.22 / -7.85$ on BLG41 and $-3.22 / -7.85$ on BLG01 — while the faint trio does not: m18_d17 reaches 45 and 80 mas decade RMS on only ~2400 usable epochs. It is a poor star.

m18_d04 is the neighbour blended with the target, 4.5 pixels away under ~7 pixel seeing. It is a legitimate member of the calibration set — the target at $I = 18.13$ is fainter than the $I < 18$ companion threshold, so it does not trigger rejection — but its centroid tracks the target's brightening through the event, and it is not an independent object.

Coordinates. Pixel positions are per field, since the cut-outs are offset by about 1.1 pixels; sky positions come from inverting each field's own map and the two fields agree on them to 2 to 9 mas.

star	I	BLG41 pixel	BLG01 pixel	RA (J2000)	Dec (J2000)
target	18.13	150.4, 150.4	151.6, 150.2	17:52:38.085	-31:47:36.21
m18_d04	18.14	146.5, 148.3	147.6, 148.0	17:52:38.210	-31:47:37.06
m18_d17	18.09	134.5, 143.6	135.6, 143.4	17:52:38.590	-31:47:39.00
m18_d48	18.05	198.6, 145.4	199.8, 145.0	17:52:36.585	-31:47:38.10
m16_d24	16.68	161.3, 172.2	162.4, 172.1	17:52:37.746	-31:47:27.53
m16_d30	16.63	120.4, 154.6	121.4, 154.5	17:52:39.023	-31:47:34.60
m16_d31	16.75	138.7, 178.7	139.8, 178.6	17:52:38.436	-31:47:24.98

7. Chi-squared of the monthly binned residuals

For every star the residuals are binned in sidereal months; each bin contributes $(\text{mean}/\text{standard error})^2$ in both axes, and $\text{DoF} = 2 \cdot N_{\text{bins}} - 4$. Values are not divided by DoF.

	BLG41	BLG01
stars measured	260	235
median DoF	182	182
target χ^2	1500 (DoF 180)	1313 (DoF 180)
target χ^2/DoF	8.3	7.3
field median χ^2/DoF	6.5	6.8

The target is unremarkable. The field median of χ^2/DoF near 7 is itself informative: the standard error of the mean within a month **understates the true bin uncertainty by about $\sqrt{7} \approx 2.6$** , because the noise within a month is correlated rather than independent. The effect is multiplicative and nearly independent of magnitude, and the faintest stars come closest to unity, since for them genuine photon noise — which *is* independent between exposures — dominates.

What the error bars in the figures are. Each plotted point is the mean position of the frames in one sidereal-month bin. Its *standard error* (s.e.) is $\sigma/\sqrt{N}$, with σ the scatter of the individual frames inside the bin and N their number — the uncertainty of the mean, not the spread of the frames. For the target a typical bin has $\sigma \approx 21\text{--}23$ mas and $N \approx 75$, hence an s.e. of about 2.5 mas, which is what the bar would be if the frames were independent. They are not, so the

bars drawn are the s.e. multiplied by an empirical factor, $\sqrt{(\chi^2/\text{DoF})}$ of that star's bin means about zero, computed per star and per axis and quoted in each panel title (2.5–3 for the target, 1.6–2 for the best-measured comparison stars): the standard error the frame count suggests, inflated to match the scatter the bins actually show.

8. A residual systematic that survives every correction

Fitting a straight line to each star's residuals **within each season separately** leaves slopes that are real and unexplained.

	BLG41	BLG01
median per-star season-slope rms	2.90 / 2.80 mas	2.93 / 2.90 mas
expected from white noise	0.99 mas	0.88 mas
excess over white noise	2.9x	3.3x

They are **not noise**. Matching 181 stars between the fields through OGLE and correlating each star's ten-season slope pattern gives a median correlation of **+0.795** in one axis and **+0.754** in the other, with **96 to 97% of stars positive**. Two independent reductions, independent frames, essentially every star showing the same season-by-season pattern: about 79% of the slope variance is shared systematic rather than measurement noise. The target is entirely typical in this respect, at the 51st to 75th percentile among stars within 0.4 mag of it.

The cause has not been identified. Nine candidates were tested and eliminated:

candidate	test
DCR / colour	per-season slope against colour
DCR colour resolution	full refit with 12 and with 20 colour bins instead of 6
calibrator contamination	second clipping iteration on the season scale, refit without the 21 + 22 stars with 2σ -high season-offset χ^2
season-edge epochs	refit without the first and last 14 nights of every season (5% of the epochs)
high-airmass epochs	refit without $\sec z > 1.3$ (30% of the epochs)
blending	against nearest-neighbour distance
field distortion	spatial correlation; quadratic in position
pixel phase	slope against within-season phase drift
SysRem cadence	refit per season instead of per decade

The colour-bin refits are the direct test of the DCR model. The 18 DCR parameters are fitted independently in each colour bin, so with 12 or 20 bins (19 or 11 stars per bin instead of 37) any colour dependence too fine for six bins would have been absorbed. Nothing changed:

	6 bins (v8)	12 bins	20 bins
slope rms, BLG41	2.90 / 2.80	2.98 / 2.77	2.92 / 2.80
slope rms, BLG01	2.93 / 2.90	2.97 / 2.89	2.92 / 2.83
cross-field correlation	+0.795 / +0.754	+0.794 / +0.742	+0.796 / +0.766

bright-star RMS, BLG41	4.80 / 4.83	4.74 / 4.83	4.74 / 4.83
bright-star RMS, BLG01	4.75 / 4.95	4.78 / 4.95	4.73 / 4.96
target RMS, BLG41	23.77 / 19.67	23.80 / 19.72	23.78 / 19.66
target RMS, BLG01	21.51 / 20.45	21.47 / 20.40	21.53 / 20.34
target PM, BLG41	−1.966 / −7.019	−1.879 / −7.020	−1.909 / −7.116
target PM, BLG01	−1.514 / −7.073	−1.684 / −7.048	−1.595 / −7.105

All in mas or mas/yr, X / Y. The cross-field correlation, which is the quantity that measures the shared systematic, is unchanged to the third decimal, the bright-star RMS moves by at most 0.06 mas and the target's proper motion by at most 0.17 mas/yr, or 0.75 of its error. Six colour bins already capture all the colour dependence there is; the six-bin solution is kept.

The season-scale clipping asks whether the frame is being pulled by calibrators that themselves carry coherent season offsets — the decade RMS used for the v8 cut averages such offsets away (a 16.75-mag star among the comparison set has its X residual at +15 mas for two seasons and −15 mas for one, with a perfectly normal decade RMS). Applying a second 2σ cut, in log space per magnitude bin, on each star's season-offset χ^2/DoF — $\Sigma(\text{season mean} / \text{s.e.})^2$ over seasons and axes — removes 21 stars in BLG41 and 22 in BLG01, two-thirds of them at 16–17 mag; the target sits at $z = +0.9 / +0.6$, inside the population. Refitting on the 239 + 213 survivors changes nothing on the stars common to both solutions:

common stars	BLG41	BLG01
RMS, $I < 17$ (117 / 105 stars)	4.66 / 4.77 → 4.65 / 4.80	4.61 / 4.95 → 4.60 / 4.94
RMS, 17–18	+1.0% / +0.5%	−1.5% / +1.5%
RMS, 18–19	+0.1% / +2.4%	−0.9% / +1.1%
Gaia scatter, RUWE-clean (181 / 146)	0.967 / 0.558 → 0.934 / 0.572	1.067 / 0.629 → 0.982 / 0.661
target PM	−1.966 / −7.019 → −2.003 / −6.997	−1.514 / −7.073 → −1.770 / −7.024

Per-star RMS moves by under 1% with half the stars going each way; the Gaia scatter improves in X and worsens in Y; the target's proper motion moves by 0.04 and 0.26 mas/yr (0.1σ and 1.1σ), bringing the two fields to within 0.23 mas/yr in X. The apparent 3% gain in the quoted bright-star RMS (4.64 / 4.72 and 4.58 / 4.94) is entirely the removal of the worst stars from the median, not an improvement of the solution. The per-epoch frame with some 240 stars is insensitive to these 8–9%, so calibrator contamination is not the origin of the slopes, and v8 is kept as the reference.

The two epoch cuts ask whether the slopes are made at the season edges or at high airmass. They are nearly independent selections: the first and last 14 nights of a season hold only 25–50 epochs each at a median sec z of 1.33–1.40, but frames at sec $z > 1.3$ are taken at the start and end of every night all season long, so only 12% of them fall in the edge trim. Both were refitted with the v8 stars and fit otherwise unchanged (the truncated 2026 season keeps its final, mid-season nights). Compared on exactly the surviving epochs, so that dropping the noisiest frames is not mistaken for a better solution:

on the common epochs	edge trim, BLG41 / BLG01	sec $z \leq 1.3$, BLG41 / BLG01
epochs kept	96% / 96%	78% / 77%

RMS change, $I < 17$	+0.6 / +0.1%, -0.1 / -0.4%	-0.6 / -0.2%, 0.0 / +0.1%
RMS change, 17–18	0.0 / -0.4%, +0.4 / +0.1%	-0.3 / -3.0%, -3.9 / -3.9%
RMS change, 18–19	+0.6 / -0.5%, +1.1 / +0.2%	-0.5 / -1.0%, -1.0 / -1.0%
RMS change, target	-0.1 / -0.5%, +2.1 / +1.1%	-2.6 / 0.0%, -3.9 / +2.0%
target PM	-1.926 / -6.988, -1.583 / -7.056	-1.530 / -7.195, -1.583 / -7.187
Gaia scatter	0.985 / 0.598, 1.039 / 0.648	0.995 / 0.596, 1.027 / 0.698
slope rms	2.68 / 2.56, 2.69 / 2.67	2.73 / 2.79, 2.81 / 2.72
excess over white noise	2.9 / 3.2	2.6 / 2.9
cross-field correlation	+0.770 / +0.734	+0.770 / +0.730

The edge trim is a null: on the common epochs nothing changes, and the 7% lower slope rms is only the shorter season. The airmass cut is a small, real improvement of the fit for stars fainter than $I = 17$ and for the target's X scatter — the differential-refraction and high-airmass PSF model evidently does worst on faint stars — and it moves the target's proper motion in BLG41 by 0.44 mas/yr in X (1.3σ), bringing the two fields to within 0.05 mas/yr; that is one number moving within its error, at the cost of 30% of the data. Either way the slope excess and the cross-field correlation survive with three-quarters of the epochs at $\sec z < 1.3$, so the shared systematic is made neither at the season edges nor at high airmass.

The per-season SysRem test is the most informative. SysRem represents residuals as a per-source coefficient times a per-epoch mode, which is exactly the structure the slopes appear to have; giving it its own two components in each season — ten times the freedom — should have absorbed them. Instead the slopes grew and every other metric degraded by 5 to 11%, because each season's correction is constrained by about 1900 epochs instead of 19 000 and only 88 to 90% of epochs received any correction at all. The slopes are therefore **not** a low-rank common mode within seasons, consistent with their being per-star and spatially uncorrelated.

Since the effect reproduces between two cut-outs of the same stars on the same nights, it is in the images rather than in the reduction. It is the likely origin of the χ^2/DoF of about 7 in section 7, of the factor 2.6 by which the internal proper-motion errors are optimistic, and of the floor on what this dataset can deliver.

9. Overlap of the two fields

Stars are identified through the OGLE catalogue; matching the two source lists directly is ambiguous because a close pair may be split in one field and merged in the other.

	BLG41	BLG01
sources	594	621
matched to an OGLE entry	551 (93%)	497 (80%)
distinct OGLE stars behind them	454	394

The gap between the last two rows is the blending, measured: 551 detections correspond to only 454 stars on BLG41, so about 97 detections are duplicates.

OGLE I	BLG41	BLG01	in both	distinct	found in both
< 15	28	22	22	28	78.6%

15 – 16	54	42	42	54	77.8%
16 – 17	172	149	146	175	83.4%
17 – 18	135	119	117	137	85.4%
18 – 19	55	44	41	58	70.7%
all brighter than 19	444	376	368	452	81.4%

About four stars in five are detected in both fields, flat from the brightest down to $I = 18$. **There is nothing fainter than $I = 19$ to compare:** our source lists span 12.87 to 18.98 in both fields, while OGLE itself reaches 21.40 and holds 2660 entries fainter than 19. That is the KMT detection limit in this crowding, not a catalogue cut.

10. Files

file	content
report_v8_RMS.png	residual RMS against OGLE I , per axis and field
report_v8_target_beforeafter.png	the target before and after proper-motion detrending
report_v8_pm_vs_gaia.png	our absolute proper motion against Gaia's, after the tie
report_v8_chi2.png	χ^2 distribution of the monthly binned residuals, target marked
report_v8_radec_<field>_<tag>.png	RA and Dec residuals against time and against each other, colour-coded by time; bi-dropped
report_v8_motion_<field>_<tag>.png	2x2 motion figures, proper motion retained and removed; same binning cuts as the RA/Dec
source_motion_v8_<field>[_<tag>].csv	per-epoch positions for the target and each comparison star

CSV columns are JD, X_mas, Y_mas, errX_decade, errY_decade, errX_season, errY_season, season, with positions relative to that star's own mean and **proper motion not removed**. Headers carry the star's magnitude, pixel position, sky position and absolute proper motion.

In the motion figures the columns are the two axes, **X left and Y right**; the top row keeps the proper motion and the bottom row removes it. The legend quotes the motion along the pixel axis and the heading quotes it on the sky. **Pixel X runs opposite to RA**, so those two carry opposite signs in X.

Figures




Figure 1. Our absolute proper motion against Gaia's, per axis and per field, after the affine tie with the **full six-parameter gauge** removed. Dashed line is 1:1. The scatter quoted in each panel is the robust standard deviation of the difference; it improves by a factor of three to four over a tie that removes only the constant. The target is the red star.

report_v8_pm_vs_gaia.png




Figure 2. χ^2 of the sidereal-month binned residuals for every calibration star, both axes combined, **not divided by DoF**. Each bin contributes $(\text{mean} / \text{standard error})^2$ per axis and $\text{DoF} = 2N_{\text{bins}} - 4$. The red line marks the target, which sits in the body of the distribution rather than the tail.

report_v8_chi2.png




Figure 3. **BLG41** — the target, $l = 18.13$. **Position and proper motion removed**: these are the residuals. Left: RA residual against time. Middle: Dec residual against time. Right: the two against each other with time colour-coded, where a well-behaved star shows an unordered cloud about the origin and any systematic drift would appear as an ordered progression of colour; consecutive months are joined by a thin grey line. Sidereal-month bins; error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means, quoted in each panel title, so that they carry the real bin-to-bin scatter (section 7). Bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped: the season-edge months, with a handful of high-airmass frames, otherwise dominate the plot.

report_v8_radecl_BLG41_target.png




Figure 4. **BLG41** — comparison star $l = 16.68$, 24 pix from the target — the best measured of the set. **Position and proper motion removed**: these are the residuals. Left: RA residual against time. Middle: Dec residual against time. Right: the two against each other with time colour-coded, where a well-behaved star shows an unordered cloud about the origin and any systematic drift would appear as an ordered progression of colour; consecutive months are joined by a thin grey line. Sidereal-month bins; error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means, quoted in each panel title, so that they carry the real bin-to-bin scatter (section 7). Bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped: the season-edge months, with a handful of high-airmass frames, otherwise dominate the plot.

report_v8_radecl_BLG41_m16_d24.png




Figure 5. **BLG41** — comparison star $l = 16.63$, 30 pix from the target. **Position and proper motion removed**: these are the residuals. Left: RA residual against time. Middle: Dec residual against time. Right: the two against each other with time colour-coded, where a well-behaved star shows an unordered cloud about the origin and any systematic drift would appear as an ordered progression of colour; consecutive months are joined by a thin grey line. Sidereal-month bins; error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means, quoted in each panel title, so that they carry the real bin-to-bin scatter (section 7). Bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped: the season-edge months, with a handful of high-airmass frames, otherwise dominate the plot.

report_v8_radec_BLG41_m16_d30.png




Figure 6. **BLG01** — the target, $l = 18.13$. **Position and proper motion removed**: these are the residuals. Left: RA residual against time. Middle: Dec residual against time. Right: the two against each other with time colour-coded, where a well-behaved star shows an unordered cloud about the origin and any systematic drift would appear as an ordered progression of colour; consecutive months are joined by a thin grey line. Sidereal-month bins; error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means, quoted in each panel title, so that they carry the real bin-to-bin scatter (section 7). Bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped: the season-edge months, with a handful of high-airmass frames, otherwise dominate the plot.

report_v8_radec_BLG01_target.png




Figure 7. **BLG01** — comparison star $l = 16.68$, 24 pix from the target — the best measured of the set. **Position and proper motion removed**: these are the residuals. Left: RA residual against time. Middle: Dec residual against time. Right: the two against each other with time colour-coded, where a well-behaved star shows an unordered cloud about the origin and any systematic drift would appear as an ordered progression of colour; consecutive months are joined by a thin grey line. Sidereal-month bins; error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means, quoted in each panel title, so that they carry the real bin-to-bin scatter (section 7). Bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped: the season-edge months, with a handful of high-airmass frames, otherwise dominate the plot.

report_v8_radec_BLG01_m16_d24.png




Figure 8. **BLG01** — comparison star $l = 16.63$, 30 pix from the target. **Position and proper motion removed**: these are the residuals. Left: RA residual against time. Middle: Dec residual against time. Right: the two against each other with time colour-coded, where a well-behaved star shows an unordered cloud about the origin and any systematic drift would appear as an ordered progression of colour; consecutive months are joined by a thin grey line. Sidereal-month bins; error bars are the standard error within the bin scaled by the empirical factor $\sqrt{\chi^2/\text{DoF}}$ of the bin means, quoted in each panel title, so that they carry the real bin-to-bin scatter (section 7). Bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped: the season-edge months, with a handful of high-airmass frames, otherwise dominate the plot.

report_v8_radec_BLG01_m16_d30.png




Figure 9. **BLG41** — the target, $l = 18.13$. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{\chi^2/\text{DoF}}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_target.png




Figure 10. **BLG41** — $l = 18.14$, 4.5 pix away — **blended with the target**, see section 7. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{\chi^2/\text{DoF}}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at sec $z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_m18_d04.png




Figure 11. **BLG41** — $l = 18.09$, 17 pix away — poorly measured, see section 7. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_m18_d17.png




Figure 12. **BLG41** — $l = 18.05$, 48 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_m18_d48.png




Figure 13. **BLG41** — $l = 16.68$, 24 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{\chi^2/\text{DoF}}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at sec $z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_m16_d24.png




Figure 14. **BLG41** — $I = 16.63$, 30 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_m16_d30.png




Figure 15. **BLG41** — $I = 16.75$, 31 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG41_m16_d31.png




Figure 16. **BLG01** — the target, $l = 18.13$. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at sec $z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG01_target.png




Figure 17. **BLG01** — $l = 18.14$, 4.5 pix away — **blended with the target**, see section 7. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG01_m18_d04.png




Figure 18. **BLG01** — $l = 18.09$, 17 pix away — poorly measured, see section 7. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG01_m18_d17.png




Figure 19. **BLG01** — $l = 18.05$, 48 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{\chi^2/\text{DoF}}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at sec $z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG01_m18_d48.png




Figure 20. **BLG01** — $I = 16.68$, 24 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.
 report_v8_motion_BLG01_m16_d24.png




Figure 21. **BLG01** — $I = 16.63$, 30 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at $\sec z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.
 report_v8_motion_BLG01_m16_d30.png




Figure 22. **BLG01** — $l = 16.75$, 31 pix away. Binned in one sidereal month, error bars are the standard error within the bin scaled by the empirical factor $\sqrt{(\chi^2/\text{DoF})}$ of the bin means (quoted in the panel titles, computed on the proper-motion-removed residuals and applied to both rows); bins with fewer than 10 epochs and epochs at sec $z > 1.3$ are dropped. Columns are the two axes, **X left and Y right**; the top row keeps the proper motion, the bottom row removes it. The heading gives the absolute sky motion with the full gauge removed; the legend gives it along the pixel axis, and **pixel X runs opposite to RA**, so the two carry opposite signs in X.

report_v8_motion_BLG01_m16_d31.png